

Varun Mundale

Lead Software Engineer - Distributed Systems | Backend Engineering | Financial Data Platforms

☎ +91 9987370106 📍 Bengaluru, India ✉ varunmundale@gmail.com

🌐 <https://www.linkedin.com/in/varun-mundale-6b8337a6/>

👤 PROFESSIONAL SUMMARY

Lead Software Engineer with **8+** years of experience designing and building **large-scale distributed systems**, backend services, and financial data platforms. Led **end-to-end engineering initiatives** from architecture to production, delivering measurable improvements in **latency, data freshness, and infrastructure efficiency**. Specialized in **event-driven architectures**, stream processing, and financial systems, with a strong focus on **correctness, reliability, and operational excellence**.

🧰 EXPERIENCE

• **Stripe (via Recko acquisition)** Bengaluru, India
Lead Software Engineer Sep 2021 – April 2026

Highlight: 12x faster end-to-end financial reporting (12hrs → less than 1hr) · 10-20x lower dashboard latency · Reconciliation scaled to 10K+ merchants · Incident classes caught pre-production

Led the architecture and delivery of high-impact initiatives across Stripe's **Revenue Recognition** and **Reconciliation platforms**, modernizing financial data processing and reporting infrastructure while improving scalability, performance, and data freshness. Owned **on-call and reliability** for the platform, driving a shift from reactive incident response to **systematic pre-production prevention**.

Incremental Reporting & Serving Platform

(Mar 2025 - April 2026)

- **Highlight: 3x faster reporting (3hrs → less than 1hr) · P99 latency 10-20x lower (10-20s → sub-second)**
- Led architecture and implementation of the next-generation reporting platform, evaluating **multiple architectural approaches** to balance recoverability, data freshness, query latency, and infrastructure cost.
- Designed an incremental **report materialization framework** on **Iceberg transactional upserts**, eliminating full-table recomputation.
- **Integrated Apache Pinot** as the analytical serving layer; designed replay and recovery workflows using Iceberg snapshots and time travel for reliable backfills.
- **Technologies:** Spark, Scala, Iceberg, Apache Pinot, Kafka, Airflow2

Incremental Processing Architecture

(Sept 2023 - Feb 2025)

- **Highlight: 4x faster pipeline (12hrs → 3hrs) · \$600/day cost reduction.** Org's critical reporting freshness milestone unlocked.
- **Redesigned** the Revenue Recognition pipeline from whole-world batch to an **incremental delta-based architecture**, processing only changed financial entities while preserving accounting correctness.
- Decomposed monolithic Spark workflows into change-based processing, collaborating cross-functionally to validate architecture and ensure seamless downstream integration.
- **Technologies:** Spark, Scala, Airflow, Trino, Iceberg, MongoDB CDC, Flink

Reconciliation Platform

(September 2021 - August 2023)

- **Highlight: Merchants scaled to 10K+ Post-acquisition platform build (Recko → Stripe)**
- Led the technical design and implementation of Stripe's reconciliation platform following Recko's acquisition, delivering **Alpha, Private Beta and Public Beta releases**.
- Built **configurable** reconciliation workflows supporting merchant **opt-in/opt-out**, date-based matching and compliance requirements, contributing to the successful onboarding of 10K+ Stripe merchants.
- **Technologies:** Spark, Scala, Airflow, Iceberg

On-Call & Reliability

(Cross-cutting)

- As the **Oncall DRI** identified a broader **systemic reliability problem** rather than treating each production incident as an isolated issue.
- **Consolidated and categorized historical incidents** by root cause and failure patterns. using Incidents AI - this revealed several recurring systemic gaps - merchant onboarding, missing business-level validation, low staging fidelity, undocumented critical flows.

- **Proposed and built a production-like staging subsystem** that empowered my team and dependent teams with **dual upstream modes** (production and staging) and **configurable input-data variants** to replay incident-prone edge cases end-to-end. This way I turned the **reactive reliability model into a proactive prevention**.
- Implemented **business invariant checkers** over real staging output, converting **production incident patterns into automated correctness** checks; documented critical flows to speed up on-call debugging and on-callee onboarding.

• Flipkart

Bengaluru, India

Software Engineer 2

May 2019 – August 2021

- **Highlight: DRI. Financial data platform made self-healing. Big Billion Day peak handled. Regulatory compliance made automated**
- **Data Correctness & Completeness:** As the DRI, **designed and built** a platform to proactively **detect, classify, and automatically remediate accounting discrepancies** within Flipkart's financial systems. Enabled faster issue detection and resolution while providing a unified observability platform for engineering, finance, and product teams.
- **Big Billion Day Readiness:** Contributed to **performance engineering** of critical accounting services by executing large-scale load and stress tests, identifying bottlenecks, and implementing **optimizations** to ensure **reliable operation** during **peak** shopping events.
- **Tax Compliance Framework:** Developed an **automated testing framework** for validating TDS/TCS regulatory changes, enabling timely rollout of government **compliance** updates with improved reliability and reduced manual validation effort.
- **Technologies:** Java, Python, Spark, Kafka, MongoDB, HBase, Locust

• Tesco

Bengaluru, India

Software Engineer 1

July 2018 – May 2019

- **Highlight: Order distribution backbone for Grocery Home Shopping made reliable**
- **Redesigned** the **queuing** system for the Grocery Home Shopping (GHS) Orders API, improving the reliable distribution of customer orders from the central order management system to retail store systems.
- **Technologies:** C# .NET, MySQL

• Flydubai Airlines

Hyderabad, India

Software Engineer Intern

June 2017 – August 2019

- **Highlight: Triage for email and call center data automated**
- Developed an **ML-based system** to **classify** email threads and call center summaries into 160 business-specific categories.
- **Technologies:** Python, TensorFlow

PERSONAL PROJECTS

• dbgit – Git for Database Schemas

 github.com/varunmundale/database-version-control

Personal Project

July 2026 – August 2026

- **Highlight: Git-style branch / diff / merge for database schemas · Team-safe · Multi-dialect**
- **Built** a tool that lets developers **branch a database like code**: each branch is a real, private copy forked in seconds, so schema changes can be tried, reviewed and thrown away without touching production.
- **Diff and merge** show exactly what changed down to the column and index, flag only **genuine conflicts**, and are **rehearsed on a scratch database** first; safe for multiple users working concurrently.
- **Technologies:** Java, jOOQ, JSqlParser, PostgreSQL, MySQL, H2, Docker

SKILLS

• Java • Scala • Spark • Python • Ruby • MySQL • Pinot • Trino • MongoDB • Hive • HBase • Kafka • Flink
• Airflow • Jenkins • Metabase • Kubernetes • Docker • GitHub • IntelliJ • Claude

EDUCATION

• IIIT Hyderabad

Sept 2016 – June 2018

Master of Technology in Computer Science & Engineering

Hyderabad, India

• Mumbai University (FCRIT)

August 2013 – June 2016

Bachelor of Engineering in Information Technology

Navi Mumbai, India

• Mumbai University (VESP)

June 2010 – June 2013

Diploma in Information Technology

Mumbai, India